

Formula 1 Racing Management Database for Enhanced Performance Analytics

Vishnu Krishnakumar Menon
UB ID: vk53
vk53@buffalo.edu

Sharvari Sunil Mayekar
UB ID: smayekar
smayekar@buffalo.edu

Abstract - This report presents the design and implementation of a relational database that manages Formula 1 racing data. The database is developed to support advanced performance analytics and decision-making by racing teams and analysts. It integrates information across seasons, races, drivers, teams, circuits, and various race-related events such as qualifying races, pit stops, penalties, and team standings. By enforcing data integrity and complex queries, this database system overcomes the limitations of excel-based solutions. The design includes 10 interrelated tables, each populated with 3000 synthetic records to simulate a realistic environment. An entity-relationship (E/R) diagram describes the relationships and constraints that support the system.

I. PROBLEM STATEMENT

Formula 1 (F1) is the highest class of international open-wheel, single-seat automobile racing, widely considered the apex of motorsport. Each race season generates vast amounts of interconnected data ranging from race scheduling and driver statistics to penalty records and pit stop timings. Managing and analysing these data points in real time is critical for teams to remain competitive. The questions we aim to address include:

1. How can we effectively integrate diverse Formula 1 data (seasons, races, drivers, teams, circuits) into a single, coherent system?
2. Which database design features best support complex queries (e.g., correlating qualifying positions with race outcomes, tracking pit stop efficiencies)?
3. How can real-time updates be facilitated to ensure that race information (such as penalties or pit stops) remains current and reliable?

A relational database is necessary instead of an Excel file due to **scalability**, **data integrity**, and **real-time querying**. Excel struggles with enforcing relationships, handling high-volume data, and supporting concurrent access without risking data corruption and performance issues. A PostgreSQL-based solution offers robust foreign key constraints, transactional integrity, and advanced query capabilities crucial for analysing dynamic Formula 1 data.

In Formula 1, even marginal improvements in lap times or pit stop strategies can decide race outcomes. This competitive edge depends on processing and interpreting vast, rapidly changing datasets. Traditional spreadsheet software cannot handle complex relationships or provide real-time insights. As a result, teams risk inaccurate or outdated information when relying on spreadsheets alone. A well-designed relational database addresses these challenges by:

1. **Maintaining Data Integrity:** Ensuring that all race records, driver profiles, and team details are consistently linked.
2. **Enabling Advanced Analytics:** Allowing teams to run complex queries that connect qualifying results, race outcomes, and pit stop data in real-time.
3. **Supporting Concurrent Usage:** Multiple users such as team managers, strategists, and analysts can access and update the database simultaneously without risking data conflicts or loss.

By consolidating the data into a dedicated Formula 1 racing management database, this project offers:

1. **Real-Time Decision Support:** Up-to-date records of race activities (penalties, pit stops, driver changes) enable immediate strategy adjustments.
2. **Enhanced Performance Analysis:** Advanced SQL queries allow deeper insights into driver performance trends, pit stop efficiency, and overall team standings.
3. **Scalable Architecture:** As new seasons, circuits, and drivers are added, the database structure adapts without sacrificing integrity or speed.

This project is crucial because it directly supports the continuous pursuit of competitive advantage in Formula 1. Teams leveraging comprehensive data analytics gain a strategic edge, refining everything from car setups to pit stop strategies in a sport where fractions of a second often determine victory.

II. TARGET USER

The primary users of this Formula 1 racing management database are the race teams and their associated support assistants, while administration is handled by specialized IT and database professionals within the organization. Below is a detailed breakdown of the target users:

Primary Users:

1. **Race Team Managers and Strategists:** These individuals are responsible for making real-time decisions during race events. They use the database to monitor driver performance, analyze race results, and develop strategies based on data insights.
2. **Performance Analysts and Data Scientists:** Analysts dive deep into historical and real-time data to identify performance trends and generate predictive models. They are tasked with developing detailed reports and dashboards for long-term strategic planning.
3. **Engineers and Technical Staff:** Engineers use the database to understand the interplay between

various technical factors such as pit stop time, car performance metrics and penalties.

Administrative Users:

- Database Administrators:** They are responsible for the deployment, maintenance, and security of the PostgreSQL database. They ensure the system is always up to date, perform regular backups, manage user permissions, and optimize performance for complex queries.
- IT Support:** These professionals manage the integration of the database with other systems, such as third-party data sources (e.g. weather updates), and analytics platforms. They ensure seamless data flow and system operability.

Consider a top-tier Formula 1 team (Ferrari) preparing for an upcoming season. On race weekends, team managers and strategists are stationed in a dedicated operations room equipped with large displays showing real-time dashboards. These dashboards are built on top of the PostgreSQL database, which integrates data from multiple sources—including live telemetry from race cars, pit stop times, and historical performance data.

As the race unfolds, performance analysts monitor live queries that track qualifying results, pit stop timings, and penalty records. Based on these insights, strategists make rapid decisions—such as adjusting pit stop timings or managing driver's performance to maximize the team's performance. Behind the scenes, DBAs and IT support ensure that the data remains consistent, secure, and quickly accessible to all users.

This integrated approach not only enhances the team's in-race performance but also provides valuable historical data for long-term improvements. The combination of detailed real-time analytics and a robust data management system transforms raw race data into a strategic asset, demonstrating the critical role of a dedicated database over traditional tools like Excel.

III. ENTITY-RELATIONSHIP (E/R) DIAGRAM

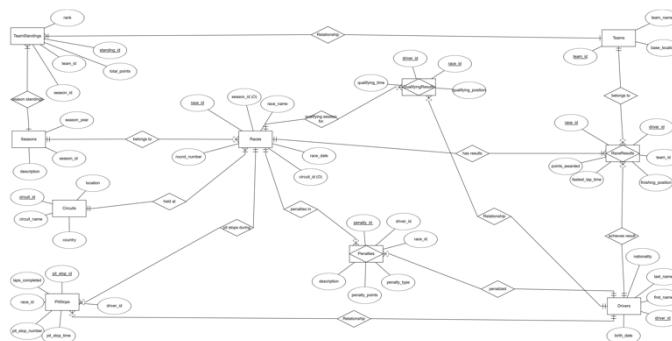


Figure 1: Formula 1 E/R diagram

The ER diagram for the Formula 1 Racing Management Database consists of ten interrelated entities, each designed to capture a specific aspect of the racing domain while ensuring data integrity through primary keys, composite keys, and foreign key constraints. Link has been provide for enhanced view of the diagram (<https://postimg.cc/3dsrNBm5>). The diagram starts with the Seasons entity, which stores information about each racing season. It includes a unique

primary key, `season_id`, along with attributes such as `season_year` and `description`. Each season can host multiple races, creating a one-to-many relationship with the Races entity. The Races entity represents individual race events and has its own primary key, `race_id`. It also includes foreign keys, such as `season_id` and `circuit_id`, which associate each race with a specific season and a race track, respectively, as recorded in the Circuits entity. The Circuits entity uses `circuit_id` as its primary key and contains details like `circuit_name`, `location`, and `country`. The RaceResults entity captures the outcomes of races. It has a composite primary key comprised of `race_id` and `driver_id`, meaning that the unique identification of a result depends on the combination of the race and the driver. This entity also includes a foreign key, `team_id`, which links each result to its corresponding team from the Teams entity. The Teams entity is identified by its primary key, `team_id`, and consists of attributes such as `team_name` and `base_location`. It is directly associated with both race results and team performance summaries. Drivers are stored in the Drivers entity, which uses `driver_id` as the primary key and includes attributes such as `first_name`, `last_name`, `nationality`, and `birth_date`. Each driver can participate in multiple races and, therefore, can appear in several other entities. For example, the QualifyingResults entity captures qualifying session details for each driver in a race and employs a composite primary key of `race_id` and `driver_id`, linking back to both Races and Drivers.

Additionally, the PitStops entity logs individual pit stop events. It has a unique primary key, `pit_stop_id`, along with foreign keys for `race_id` and `driver_id`, indicating which race and driver the pit stop applies to. The Penalties entity records any penalties incurred by drivers, featuring a unique primary key, `penalty_id`, and foreign keys for `race_id` and `driver_id` to specify the race and the driver involved. The TeamStandings entity summarizes the overall performance of teams across seasons. It features a unique primary key, `standing_id`, and includes foreign keys, `season_id` and `team_id`, to associate each standing with a specific season and team.

Throughout the diagram, the cardinalities are primarily one-to-many; for instance, one season can host multiple races, one race can generate multiple race results and qualifying results, and one driver can be associated with several pit stops or penalties. The use of composite keys in entities like RaceResults and QualifyingResults ensures that the combination of race and driver uniquely identifies each record. Meanwhile, foreign key constraints maintain referential integrity by linking child records to their parent entities. This comprehensive design provides a clear and robust framework for managing the complex data relationships inherent in Formula 1 racing.

IV. DATASET ACQUISITION

A fake dataset for the Formula One Racing Management Database was created using a Python script. This script utilizes the Faker library to generate realistic names, dates, and locations, and employs Python's built-in random module to produce random numeric values. The output is saved as CSV files, one for each table in the database which can be easily imported into PostgreSQL. The script starts by defining the number of records for each table. For example, it generates 20 seasons representing the years 2001

to 2020, while more data-intensive tables like Drivers, Teams, and Circuits each produce 3,000 records. This ensures that the dataset meets the minimum requirements and provides a sufficient volume for testing and advanced queries.

The following code snippet was used for producing the fake dataset. In the line – with open('seasons.csv', 'w', newline=") as file;, ‘seasons’ keyword was updated each time we generated a new table. All the 10 tables were generated using this python script by making the necessary changes(attribute names).

```
from faker import Faker
import random, csv

fake = Faker()

num_seasons = 20
num_drivers = 3000
num_teams = 3000
num_circuits = 3000
num_races = 3000
num_race_results = 3000
num_qualifying_results = 3000
num_pit_stops = 3000
num_penalties = 3000
num_team_standings = 3000

with open('seasons.csv', 'w', newline='') as file:
    writer = csv.writer(file)
    writer.writerow(['season_id', 'season_year', 'description'])
    for season in range(1, num_seasons + 1):
        season_year = 2000 + season
        description = f"Season {season_year} details"
        writer.writerow([season, season_year, description])
```

V. BOYCE-CODD NORMAL FORM

Listing functional dependencies for each relations and ensuring BCNF.

Seasons (*season_id* → *season_year*, *description*)

candidate key here is *season_id*. Since *season_id* is unique and determines the other attributes, the dependency is in BCNF. Thus, Seasons is already in BCNF, and no decomposition is needed.

Drivers (*driver_id* → *first_name*, *last_name*, *nationality*, *birth_date*)

In the Drivers relation, *driver_id* serves as the primary key. Additionally, the combination of (*first_name*, *last_name*, *birth_date*) can function as a candidate key. Since the left-hand side of the dependency is a candidate key (*driver_id*), the relation is in Boyce-Codd Normal Form (BCNF). Thus, Drivers is in BCNF.

Teams (*team_id* → *team_name*, *base_location*)

Here *team_id* is the candidate key at it is on LHS. Therefore, Teams is in BCNF.

Circuits (*circuit_id* → *circuit_name*, *location*, *country*, *circuit_name* → *circuit_id*, *location*, *country*)

Both *circuit_id* and *circuit_name* can serve as candidate keys. Since every dependency has a candidate key on the LHS, the relation is in BCNF.

Races (*race_id* → *season_id*, *race_name*, *circuit_id*, *race_date*, *round_number*)

The candidate key is *race_id*, which determines all other attributes in the Races relation, thereby meeting the BCNF condition. Thus, the relation Races is in BCNF.

RaceResults ((*race_id*, *driver_id*) → *team_id*, *finishing_position*, *points_awarded*, *fastest_lap_time*)

The composite key (*race_id*, *driver_id*) serves as the candidate key for RaceResults. Because this composite key is the only determinant and it uniquely identifies each record, the relation is in BCNF. Thus RaceResults is in BCNF.

QualifyingResults ((*race_id*, *driver_id*) → *qualifying_position*, *qualifying_time*)

The composite key (*race_id*, *driver_id*) uniquely identifies each record in QualifyingResults. Since the LHS of the dependency is a candidate key, the relation is in BCNF. Therefore, QualifyingResults is in BCNF

PitStops (*pitstop_id* → *race_id*, *driver_id*, *pitstop_number*, *pitstop_time*, *laps_completed*)

pitstop_id is the primary key for PitStops. Since it uniquely determines all other attributes, the dependency is in BCNF. Thus, PitStops is in BCNF.

Penalties (*penalty_id* → *race_id*, *driver_id*, *penalty_type*, *penalty_points*, *description*)

Penalty_id is the candidate key, and it uniquely determines the other attributes. Therefore, this dependency meets BCNF. Penalties is in BCNF.

TeamStandings (*team_standing_id* → *season_id*, *team_id*, *total_points*, *rank*)

The primary key can be either *team_standing_id* or the composite key (*season_id*, *team_id*). In both cases, the left-hand side of the dependency is a candidate key, indicating that the relation is in BCNF. Therefore , TeamStandings is in BCNF.

VI. RELATIONS AND THEIR ATTRIBUTES

Below is a list of the relations (tables) in the Formula 1 Racing Season Management database along with descriptions of their attributes, primary keys, foreign keys, default values, and actions.

1. Seasons relation:

Attributes:

season_id:

Datatype: SERIAL

Purpose: A unique identifier for each season.

Constraints: Primary Key, cannot be null, auto-generated.

season_year:

Datatype: INT

Purpose: Represents the year of the season.

Constraint: NOT NULL, Unique to prevent duplicate season entries, no default value.

description:

Datatype: TEXT

Purpose: Optional details about the season.

Constraint: Can be null.

Foreign Key Actions: This table is a master table. Other tables (e.g., Races) reference its *season_id* using ON DELETE RESTRICT which prevents the deletion of a season if related race records exist.

2. Drivers relation:

Attributes:

driver_id:

Datatype: SERIAL

Purpose: Unique identifier for each driver.

Constraint: Primary Key; not null.

first_name:

Datatype: VARCHAR(100)

Purpose: Driver's first name.

Constraint: NOT NULL.

last_name:

Datatype: VARCHAR(100)

Purpose: Driver's last name.

Constraint: NOT NULL.

nationality:

Datatype: VARCHAR(100)

Purpose: Driver's nationality.

Constraint: Can be null.

birth_date:

Datatype: DATE

Purpose: Driver's date of birth.
Constraint: Can be null.
Unique Constraint: (first_name, last_name, birth_date) ensures that duplicate driver entries are not created.
Foreign Key Actions: As an independent master table, it has no foreign keys. Other tables (RaceResults, QualifyingResults, PitStops, Penalties) reference driver_id with ON DELETE CASCADE to automatically remove dependent records if a driver is deleted.

3. Teams relation:

Attributes:

team_id:

Datatype: SERIAL

Purpose: Unique identifier for each team.
Constraint: Primary Key; auto-generated.

team_name:

Datatype: VARCHAR(100)

Purpose: The name of the team.
Constraint: NOT NULL; Unique to prevent duplicates.

base_location:

Datatype: VARCHAR(100)

Purpose: The headquarters of the team.
Constraint: Can be null.

Foreign Key Actions: This table is also independent. Other tables reference team_id (e.g. in RaceResults and TeamStandings) with ON DELETE CASCADE to ensure that when a team is removed, its associated records are automatically deleted.

4. Circuits relation:

Attributes:

circuit_id:

Datatype: SERIAL

Purpose: Unique identifier for each racing circuit.
Constraint: Primary Key; auto-generated.

circuit_name:

Datatype: VARCHAR(100)

Purpose: Name of the circuit.
Constraint: NOT NULL; Unique.

location:

Datatype: VARCHAR(100)

Purpose: City where the circuit is located.
Constraint: Can be null.

country:

Datatype: VARCHAR(100)

Purpose: Country where the circuit is located.
Constraint: Can be null.

Foreign Key Actions: As a master table, it is referenced by Races (using ON DELETE RESTRICT), preventing deletion if a circuit is in use.

5. Races relation:

Attributes:

race_id:

Datatype: SERIAL

Purpose: Unique identifier for each race event.
Constraint: Primary Key; auto-generated.

season_id:

Datatype: INT

Purpose: Links the race to a specific season.
Constraint: NOT NULL; Foreign key referencing Seasons(season_id) with ON DELETE RESTRICT.

race_name:

Datatype: VARCHAR(100)

Purpose: Name of the race event (e.g. "Grand Prix of Monaco").
Constraint: NOT NULL.

circuit_id:

Datatype: INT

Purpose: Associates the race with a particular circuit.
Constraint: NOT NULL; Foreign key referencing Circuits(circuit_id) with ON DELETE RESTRICT.

race_date:

Datatype: DATE

Purpose: Date when the race was held.
Constraint: NOT NULL.

round_number:

Datatype: INT

Purpose: The sequential order of the race in the season.
Constraint: NOT NULL.

Foreign Key Actions: Both season_id and circuit_id use ON DELETE RESTRICT to protect the integrity of race records ensuring that seasons or circuits with dependent races are not deleted.

6. RaceResults relation:

Attributes:

race_id:

Datatype: INT

Purpose: Identifier for the race event.
Constraint: Part of the composite primary key; NOT NULL; Foreign key referencing Races(race_id) with ON DELETE CASCADE.

driver_id:

Datatype: INT

Purpose: Identifier for the driver participating in the race.
Constraint: Part of the composite primary key; NOT NULL; Foreign key referencing Drivers(driver_id) with ON DELETE CASCADE.

team_id:

Datatype: INT

Purpose: Identifier for the team the driver represents during the race.
Constraint: NOT NULL; Foreign key referencing Teams(team_id) with ON DELETE CASCADE.

finishing_position:

Datatype: INT

Purpose: The final finishing position of the driver in the race.
Constraint: NOT NULL; Must be greater than 0.

Points_awarded:

Datatype: DECIMAL(5,2)

Purpose: Points awarded to the driver for the race finish.
Constraint: NOT NULL; Must be non-negative.

fastest_lap_time:

Datatype: TIME

Purpose: The fastest lap time recorded by the driver.
Constraint: Can be null.

Primary Key: Composite key (race_id, driver_id)

Foreign Key Actions: ON DELETE CASCADE on all foreign keys ensures that if the corresponding race, driver, or team is deleted, the related race result record is automatically deleted.

7. QualifyingResults relation:

Attributes:

race_id:

Datatype: INT

Purpose: Identifier for the race event.
Constraint: Part of the composite primary key; NOT NULL; Foreign key referencing Races(race_id) with ON DELETE CASCADE.

driver_id:

Datatype: INT

Purpose: Identifier for the driver in the qualifying session.

Constraint: Part of the composite primary key; NOT NULL; Foreign key referencing Drivers(driver_id) with ON DELETE CASCADE.

qualifying_position:

Datatype: INT

Purpose: The driver's finishing position in the qualifying session.

Constraint: NOT NULL; Must be greater than 0.

qualifying_time:

Datatype: TIME

Purpose: The lap time achieved by the driver during qualifying.

Constraint: NOT NULL.

Primary Key: Composite key (race_id, driver_id)

Foreign Key Actions: Both foreign keys use ON DELETE CASCADE to remove qualifying records if the corresponding race or driver is deleted.

8. Pitstops relation:

Attributes:

pitstop_id:

Datatype: SERIAL

Purpose: Unique identifier for each pit stop event.

Constraint: Primary Key; auto-generated.

race_id:

Datatype: INT

Purpose: Identifier for the race during which the pit stop occurred.

Constraint: NOT NULL; Foreign key referencing Races(race_id) with ON DELETE CASCADE.

driver_id:

Datatype: INT

Purpose: Identifier for the driver who made the pit stop.

Constraint: NOT NULL; Foreign key referencing Drivers(driver_id) with ON DELETE CASCADE.

pitstop_number:

Datatype: INT

Purpose: The sequential number of the pit stop within the race.

Constraint: NOT NULL.

pitstop_time:

Datatype: TIME

Purpose: Duration of the pit stop.

Constraint: NOT NULL.

laps_completed:

Datatype: INT

Purpose: Number of laps completed before this pit stop.

Constraint: NOT NULL; Must be greater than or equal to 0.

Primary Key: pitstop_id

Foreign Key Actions: Both foreign keys use ON DELETE CASCADE so that deletion of the associated race or driver removes the corresponding pit stop record.

9. Penalties relation:

Attributes:

penalty_id:

Datatype: SERIAL

Purpose: Unique identifier for each penalty event.

Constraint: Primary Key; auto-generated.

race_id:

Datatype: INT

Purpose: Identifier for the race where the penalty was imposed.

Constraint: NOT NULL; Foreign key referencing Races(race_id) with ON DELETE CASCADE.

driver_id:

Datatype: INT

Purpose: Identifier for the driver who received the penalty.

Constraint: NOT NULL; Foreign key referencing Drivers(driver_id) with ON DELETE CASCADE.

penalty_type:

Datatype: VARCHAR(100)

Purpose: Describes the type of penalty.

Constraint: NOT NULL.

penalty_points:

Datatype: INT

Purpose: The penalty points associated.

Constraint: NOT NULL; Must be non-negative.

description:

Datatype: TEXT

Purpose: Optional information about the penalty.

Constraint: can be null.

Primary Key: penalty_id

Foreign Key Actions: Both race_id and driver_id are set to ON DELETE CASCADE, ensuring that deletion of a race or driver also removes related penalty records.

10. TeamStandings relation:

Attributes:

team_standing_id:

Datatype: SERIAL

Purpose: Unique identifier for each team standing record.

Constraint: Primary Key; auto-generated.

season_id:

Datatype: INT

Purpose: Identifier for the season.

Constraint: NOT NULL; Foreign key referencing Seasons(season_id) with ON DELETE CASCADE.

team_id:

Datatype: INT

Purpose: Identifier for the team.

Constraint: NOT NULL; Foreign key referencing Teams(team_id) with ON DELETE CASCADE.

total_points:

Datatype: DECIMAL(7,2)

Purpose: Cumulative points earned by the team during the season.

Constraint: NOT NULL; Default value is 0.

rank:

Datatype: INT

Purpose: The overall ranking of the team for the season.

Constraint: Can be null.

Unique Constraint: Unique combination of (season_id, team_id) to ensure each team appears only once per season.

Primary Key: team_standing_id

Foreign Key Actions: Both foreign keys (season_id and team_id) are set to ON DELETE CASCADE—if a season or team is removed, the corresponding team standings record is automatically deleted.

VII. CREATING TABLES AND IMPORTING DATA

We implemented the entire schema directly in PostgreSQL using standard DDL statements: each table (Seasons, Drivers, Teams, Circuits, Races, RaceResults, QualifyingResults, PitStops, Penalties and TeamStandings) was created with CREATE TABLE, appropriate PRIMARY KEY and FOREIGN KEY constraints, UNIQUE and CHECK validations, and SERIAL columns for auto-incrementing IDs. This approach ensured referential integrity, enforced data validity, and laid the foundation for all of our queries and transactions. The screenshots of the commands are attached below.

```

Query Query History
1 -- 1. Sessions
2 v CREATE TABLE Sessions (
3 season_id SERIAL PRIMARY KEY,
4 season_year INT NOT NULL UNIQUE,
5 description TEXT
6 );
7
8 -- 2. Drivers
9 v CREATE TABLE Drivers (
10 driver_id SERIAL PRIMARY KEY,
11 first_name VARCHAR(100) NOT NULL,
12 last_name VARCHAR(100) NOT NULL,
13 nationality VARCHAR(100),
14 birth_date DATE
15 );
16
17 -- 3. Teams
18 v CREATE TABLE Teams (
19 team_id SERIAL PRIMARY KEY,
20 team_name VARCHAR(100) NOT NULL UNIQUE,
21 base_location VARCHAR(100)
22 );
23
24 -- 4. Circuits
25 v CREATE TABLE Circuits (
26 circuit_id SERIAL PRIMARY KEY,
27 circuit_name VARCHAR(100) NOT NULL UNIQUE,
28 location VARCHAR(100),
29 country VARCHAR(100)
30 );
31
32 -- 5. Races
33 v CREATE TABLE Races (
34 race_id SERIAL PRIMARY KEY,
35 season_id INT NOT NULL,
36 race_name VARCHAR(100) NOT NULL,
37 circuit_id INT NOT NULL,
38 race_date DATE NOT NULL,
39 round_number INT,
40 FOREIGN KEY (season_id) REFERENCES Seasons(season_id) ON DELETE RESTRICT
41 );
42
43 v CREATE TABLE RaceResults (
44 race_id INT NOT NULL,
45 driver_id INT NOT NULL,
46 team_id INT NOT NULL,
47 finishing_position INT NOT NULL CHECK (finishing_position > 0),
48 points_awarded DECIMAL(5,2) NOT NULL CHECK (points_awarded >= 0),
49 fastest_lap_time TIME,
50 PRIMARY KEY (race_id, driver_id),
51 FOREIGN KEY (race_id) REFERENCES Races(race_id) ON DELETE CASCADE,
52 FOREIGN KEY (driver_id) REFERENCES Drivers(driver_id) ON DELETE CASCADE,
53 FOREIGN KEY (team_id) REFERENCES Teams(team_id) ON DELETE CASCADE
54 );
55
56 -- 6. QualifyingResults
57 v CREATE TABLE QualifyingResults (
58 race_id INT NOT NULL,
59 driver_id INT NOT NULL,
60 qualifying_position INT NOT NULL CHECK (qualifying_position > 0),
61 qualifying_lap_time TIME NOT NULL,
62 PRIMARY KEY (race_id, driver_id),
63 FOREIGN KEY (race_id) REFERENCES Races(race_id) ON DELETE CASCADE,
64 FOREIGN KEY (driver_id) REFERENCES Drivers(driver_id) ON DELETE CASCADE
65 );
66
67 -- 7. PitStops
68 v CREATE TABLE PitStops (
69 pit_stop_id SERIAL PRIMARY KEY,
70 race_id INT NOT NULL,
71 driver_id INT NOT NULL,
72 pit_stop_number INT NOT NULL CHECK (pit_stop_number > 0),
73 pit_stop_time TIME NOT NULL,
74 laps_completed INT NOT NULL CHECK (laps_completed >= 0),
75 FOREIGN KEY (race_id) REFERENCES Races(race_id) ON DELETE CASCADE,
76 FOREIGN KEY (driver_id) REFERENCES Drivers(driver_id) ON DELETE CASCADE
77 );
78
79 -- 8. Penalties
80 v CREATE TABLE Penalties (
81 penalty_id SERIAL PRIMARY KEY,
82 );
83
84 -- 9. Penalties
85 v CREATE TABLE Penalties (
86 penalty_id SERIAL PRIMARY KEY,
87 race_id INT NOT NULL,
88 driver_id INT NOT NULL,
89 penalty_type VARCHAR(50) NOT NULL,
90 penalty_points INT NOT NULL CHECK (penalty_points >= 0),
91 description TEXT,
92 FOREIGN KEY (race_id) REFERENCES Races(race_id) ON DELETE CASCADE,
93 FOREIGN KEY (driver_id) REFERENCES Drivers(driver_id) ON DELETE CASCADE
94 );
95
96 -- 10. TeamStandings
97 v CREATE TABLE TeamStandings (
98 standing_id SERIAL PRIMARY KEY,
99 season_id INT NOT NULL,
100 team_id INT NOT NULL,
101 total_points DECIMAL(7,2) NOT NULL DEFAULT 0,
102 rank,
103 FOREIGN KEY (season_id) REFERENCES Seasons(season_id) ON DELETE CASCADE,
104 FOREIGN KEY (team_id) REFERENCES Teams(team_id) ON DELETE CASCADE,
105 UNIQUE (season_id, team_id),
106 );

```

After defining the schema, we loaded each CSV file into its matching table using pgAdmin's Import tool—selecting the file, mapping columns, enabling the header row, and specifying the comma delimiter—so that all our season, driver, team, circuit, race, result, penalty, pit-stop, qualifying, and standings data were imported directly and correctly into PostgreSQL.

VIII. TASK 5 – 10 SQL QUERIES

1. This query adds a new driver (ID 3001, Alex Thompson, British, born 1995-07-12) to the drivers table and then immediately retrieves that row to confirm the insert.

The screenshot shows the pgAdmin interface with two panes. The left pane contains the SQL query:

```

1 v INSERT INTO Drivers (driver_id, first_name, last_name, nationality, birth_date)
2 VALUES (3001, 'Alex', 'Thompson', 'British', '1995-07-12');
3
4 v SELECT *
5   FROM Drivers
6 WHERE driver_id = 3001;
7

```

The right pane shows the Data Output tab with the results of the query:

driver_id	first_name	last_name	nationality	birth_date
3001	Alex	Thompson	British	1995-07-12

2. This query updates the base_location of the team with team_id = 5 to "Milton Keynes" and then selects that row to confirm the change.

The screenshot shows the pgAdmin interface with two panes. The left pane contains the SQL query:

```

1 -- updating a team's base location
2 v UPDATE Teams
3   SET base_location = 'Milton Keynes'
4 WHERE team_id = 5;
5
6 v SELECT *
7   FROM Teams
8 WHERE team_id = 5;
9

```

The right pane shows the Data Output tab with the results of the query:

team_id	team_name	base_location
5	Evans Inc Racing	Milton Keynes

3. This query deletes the penalty record with penalty_id = 42 and then selects it, confirming that no row is returned.

The screenshot shows the pgAdmin interface with two panes. The left pane contains the SQL query:

```

1 -- (4) DELETE a penalty entry
2 v DELETE FROM Penalties
3 WHERE penalty_id = 42;
4
5 v SELECT *
6   FROM Penalties
7 WHERE penalty_id = 42;
8

```

The right pane shows the Data Output tab with the results of the query:

penalty_id	race_id	driver_id	penalty_type	penalty_points	description

4. This query retrieves the five most recent race winners—joining RaceResults to Races and Drivers, filtering for finishing_position = 1, ordering by race_date descending, and limiting to the top 5—outputting each race name, winner's full name, and date.

The screenshot shows the pgAdmin interface with two panes. The left pane contains the SQL query:

```

1 -- Top 5 race winners | with driver name and race date, newest first
2 v SELECT rr.race_name,
3        d.first_name || ' ' || d.last_name AS driver,
4        r.race_date
5   FROM RaceResults rr
6   JOIN Races r ON rr.race_id = r.race_id
7   JOIN Drivers d ON rr.driver_id = d.driver_id
8 WHERE rr.finishing_position = 1
9 ORDER BY r.race_date DESC
10 LIMIT 5;
11

```

The right pane shows the Data Output tab with the results of the query:

race_name	driver	race_date
Grand Prix New Amyhaven	Shelby Wilson	2025-03-10
Grand Prix Lake Marktown	Michele Caldwell	2025-03-10
Grand Prix Lake Lawrence	Kevin Harris	2025-03-06
Grand Prix West Rhondachester	Matthew Le	2025-02-27
Grand Prix Rhondafort	Jesse Reyes	2025-02-17

5. This query aggregates each driver's total points from RaceResults, orders them descending, and returns the top 10 drivers with their IDs, full names, and summed points.

```

Query  Query History
1 -- Total points per driver
2 v SELECT
3   d.driver_id,
4   d.first_name || ' ' || d.last_name AS driver,
5   SUM(rr.points_awarded) AS total_points
6 FROM RaceResults rr
7 JOIN Drivers d USING (driver_id)
8 GROUP BY d.driver_id, driver
9 ORDER BY total_points DESC
10 LIMIT 10;

```

Data Output Messages Notifications

Showing rows: 1 to 10

	driver_id	driver	total_points
1	2773	Karen Holloway	88.22
2	2310	Monique Brown	82.75
3	425	Dennis Wilson	80.09
4	2812	Jesse Reyes	79.90
5	1461	Andrew Hunt	79.50
6	1764	Charles Johnson	78.75
7	135	Elizabeth Greer	77.62
8	1071	Tracy Woodard	77.03
9	1648	Victoria Burgess	76.17
10	2965	Zachary Martinez	75.40

6. This query finds and lists all drivers whose total career points exceed the average points per driver by comparing each driver's sum of points_awarded to the overall average.

```

Query  Query History
1 -- Drivers scoring above the average total points
2 v SELECT driver, total_points
3 FROM (
4   SELECT
5     d.driver_id,
6     d.first_name || ' ' || d.last_name AS driver,
7     SUM(rr.points_awarded) AS total_points
8   FROM RaceResults rr
9   JOIN Drivers d USING (driver_id)
10  GROUP BY d.driver_id, driver
11 ) t
12 WHERE total_points > (
13   SELECT AVG(sum_pts)
14   FROM (
15     SELECT SUM(points_awarded) AS sum_pts
16     FROM RaceResults
17     GROUP BY driver_id
18   ) x
19 );

```

Query Query History

1 -- Drivers scoring above the average total points

Data Output Messages Notifications

Showing rows: 1 to 1000

	driver	total_points
1	Laura Owens	32.70
2	Suzanne Anderson	46.31
3	Timothy Smith	23.32
4	Dennis Walker	23.14
5	Jennifer Goodman	23.62
6	David Wilson	29.16
7	Sarah Jacobs	26.42
8	Vincent Farrell	32.83
9	Thomas Garcia	21.63
10	Katherine Ponce	49.11
11	Scott Evans	35.15
12	Monica Knight	26.89
13	David Faulkner	37.01

7. This query calculates for each driver, their average number of pit stops per race (by first counting stops per driver-race then averaging those counts), and returns the top five drivers with the highest average pit stops.

```

1 -- Average pit stops per driver
2 v SELECT
3   d.driver_id,
4   d.first_name || ' ' || d.last_name AS driver,
5   AVG(sub.ps_count) AS avg_pitstops
6 FROM (
7   SELECT race_id, driver_id, COUNT(*) AS ps_count
8   FROM PitStops
9   GROUP BY race_id, driver_id
10 ) sub
11 JOIN Drivers d USING (driver_id)
12 GROUP BY d.driver_id, driver
13 ORDER BY avg_pitstops DESC
14 LIMIT 5;

```

Data Output Messages Notifications

Showing

	driver_id	driver	avg_pitstops
1	2562	Amanda Eaton	1.0000000000000000
2	373	Leah Smith	1.0000000000000000
3	1933	Mary Carroll	1.0000000000000000

8. This query returns the distinct years in which “Mercer Inc Racing” appears in the team standings by joining TeamStandings with Seasons and Teams.

```

Query  Query History
1 v SELECT DISTINCT
2   s.season_year AS season_with_mercer
3   FROM TeamStandings ts
4   JOIN Seasons s ON ts.season_id = s.season_id
5   JOIN Teams t ON ts.team_id = t.team_id
6   WHERE t.team_name = 'Mercer Inc Racing';
7

```

Data Output Messages Notifications

Show

	season_with_mercer
1	2018
2	2019

9. This query assigns each result a descending row number per driver (most recent first) using a window function, then returns only the top five rows (i.e. the last five races) for each driver.

```

Query  Query History
1 -- Last 5 race results per driver
2 v SELECT *
3   FROM (
4   SELECT
5     rr.*,
6     ROW_NUMBER() OVER (PARTITION BY driver_id ORDER BY race_id DESC) AS rn
7   FROM RaceResults rr
8   ) sub
9   WHERE rn <= 5;

```

Data Output Messages Notifications

Showing rows: 1 to 1000

	race_id	driver_id	team_id	finishing_position	points_awarded	fastest_lap_time	rn
3	2653	5	2220	12	24.11	07:54:05	1
4	850	5	2769	13	3.51	14:19:31	2
5	748	5	982	10	8.33	01:59:48	3
6	565	5	2290	11	10.36	09:38:39	4
7	2662	6	938	2	5.81	03:15:57	1
8	1975	6	1393	2	7.39	23:53:12	2
9	181	7	2601	3	22.00	22:15:26	1

Total rows: 3000 Query complete 00:00:00.344

10. This function, add_race_result, takes race ID, driver ID, team ID, finishing position, points scored, and lap time, attempts to insert them into the RaceResults table, and if an entry for that same race/driver already exists (violating the unique constraint), it catches the error and issues a notice instead of failing. We defined the function below:

```

Query  Query History
1 v CREATE OR REPLACE FUNCTION add_race_result(
2   p_race INT,
3   p_driver INT,
4   p_team INT,
5   p_pos INT,
6   p_pts NUMERIC,
7   p_time TIME
8 )
9 RETURNS VOID
10 AS $$
11 BEGIN
12   INSERT INTO RaceResults
13   (race_id, driver_id, team_id, finishing_position, points_awarded, fastest_lap_time)
14   VALUES
15   (p_race, p_driver, p_team, p_pos, p_pts, p_time);
16   EXCEPTION
17   WHEN unique_violation THEN
18     RAISE NOTICE 'Result for race % and driver % already exists.', p_race, p_driver;
19   END;
20   $$ LANGUAGE plpgsql;
21
22   SELECT add_race_result(1, 3001, 5, 7, 6.00, '00:01:23');
23

```

After this, we called this function add_race_result to insert a new entry into RaceResults for race 1, driver 3001 (team 5) with a 10th-place finish, 8.5 points, and a fastest lap of 00:01:12.

```

Query   Query History
1 ✓ SELECT *
2   FROM RaceResults
3   WHERE race_id = 1
4     AND driver_id = 3001;

```

Data Output						
race_id	driver_id	team_id	finishing_position	points_awarded	fastest_lap_time	
1	3001	5	10	8.50	00:01:12	

Data Output Messages Notifications

`add_race_result`

1 add_race_result
void

Then we check if the function is working correctly by using SELECT command which retrieves the RaceResults row for race 1 and driver 3001, showing that entry's team, finishing position, points awarded, and fastest lap time.

```

1 ✓ SELECT *
2   FROM RaceResults
3   WHERE race_id = 1
4     AND driver_id = 3001;

```

Data Output						
race_id	driver_id	team_id	finishing_position	points_awarded	fastest_lap_time	
1	3001	5	10	8.50	00:01:12	

Then we call the add_race_result function again for race 1/driver 3001 with new stats, but the function's unique-violation handler catches the duplicate key and raises a notice ("Result for race 1 and driver 3001 already exists") instead of inserting a second row.

```

1 ✓ SELECT add_race_result(
2   1,          -- same race
3   3001,       -- same driver
4   5,          -- same team
5   12,         -- different position
6   15.0,        -- different points
7   '00:01:10'  -- different lap time
8 );

```

NOTICE: Result for race 1 and driver 3001 already exists.

Successfully run. Total query runtime: 102 msec.
1 rows affected.

Then we verify the data by using SELECT command again which pulls back the RaceResults row for race 1 and driver 3001—showing the existing entry (position 10, points 8.50, lap 00:01:12) and confirming that no new duplicate was added.

```

Query   Query History
1 ✓ SELECT *
2   FROM RaceResults
3   WHERE race_id = 1
4     AND driver_id = 3001;

```

Data Output						
race_id	driver_id	team_id	finishing_position	points_awarded	fastest_lap_time	
1	3001	5	10	8.50	00:01:12	

11. This procedure update_driver_nationality(p_driver_id, p_nation) looks up the row in the Drivers table with the given driver_id, sets its nationality to the provided value, and if no row was actually updated (i.e. no such driver exists), it gives a NOTICE saying "No driver with ID <id> found."

```

1 ✓ CREATE OR REPLACE PROCEDURE update_driver_nationality(
2   p_driver_id INT,
3   p_nation    VARCHAR
4 )
5 LANGUAGE plpgsql
6 AS $$
7 BEGIN
8   UPDATE Drivers
9     SET nationality = p_nation
10    WHERE driver_id = p_driver_id;
11
12 IF NOT FOUND THEN
13   RAISE NOTICE 'No driver with ID % found.', p_driver_id;
14 END IF;
15 END;
16 $$;

```

Data Output Messages Notifications

CREATE FUNCTION

Query returned successfully in 82 msec.

We verify the above function by fetching the ID, full name, and nationality for the driver with driver_id = 1, revealing that "Sara Wilson" is from "South Georgia and the South Sandwich Islands."

```

1 ✓ SELECT driver_id, first_name || ' ' || last_name AS driver, nationality
2   FROM Drivers
3   WHERE driver_id = 1;
4

```

Data Output Messages Notifications

driver_id	driver	nationality
1	Sara Wilson	South Georgia and the South Sandwich Islands

We then call the update_driver_nationality procedure to change driver 1's nationality to "Canadian," and it executes successfully.

```

1 CALL update_driver_nationality(1, 'Canadian');
2

```

Data Output Messages Notifications

CALL

Query returned successfully in 80 msec.

Then we check the driver's ID, full name, and nationality for the driver whose driver_id is 1, showing that Sara Wilson's nationality has been updated to Canadian.

```
Query  Query History
1 ✓ SELECT driver_id, first_name || ' ' || last_name AS driver, nationality
2   FROM Drivers
3   WHERE driver_id = 1;
4 |
```

Data Output Messages Notifications

driver_id	driver	nationality
1	Sara Wilson	Canadian

Then we attempt to update the nationality of the driver with ID 9999 to "Martian," but because no such driver exists, the procedure's IF NOT FOUND branch fires and raises the notice "No driver with ID 9999 found."

```
Query  Query History
1 CALL update_driver_nationality(9999, 'Martian');
2 |
```

Data Output Messages Notifications

NOTICE: No driver with ID 9999 found.
CALL

Query returned successfully in 90 msec.

IX. TASK 6 – TRANSACTION AND TRIGGERS

We set up an automatic error-logging and notification in the following way:

1. ErrorLog table: a simple audit table with an auto-incremented log_id, a timestamp (defaults to now), and a free-text err_msg column.

2. notify_error() trigger function: a PL/pgSQL routine that whenever it's called on a newly inserted row, emits a NOTICE showing the error's timestamp and message.

3. after_error_insert trigger: ties the above function to ErrorLog so that after each insert into ErrorLog, notify_error() runs and immediately prints out your logged error.

The screenshot of the trigger is attached below.

```
Query  Query History
1 -- Table to capture errors
2 ✓ CREATE TABLE ErrorLog (
3   log_id SERIAL PRIMARY KEY,
4   err_time TIMESTAMP DEFAULT CURRENT_TIMESTAMP,
5   err_msg TEXT
6 );
7
8 -- Trigger function to notify on insert
9 ✓ CREATE OR REPLACE FUNCTION notify_error() RETURNS TRIGGER AS $$
10 BEGIN
11   RAISE NOTICE 'Error logged at %: %', NEW.err_time, NEW.err_msg;
12   RETURN NEW;
13 END;
$$ LANGUAGE plpgsql;
14
15 -- 3) Trigger on ErrorLog
16 ✓ CREATE TRIGGER after_error_insert
17   AFTER INSERT ON ErrorLog
18   FOR EACH ROW
19   EXECUTE PROCEDURE notify_error();
```

Data Output Messages Notifications

CREATE TRIGGER

After this, the procedure transfer_team_points(from_team, to_team, pts) was created to encapsulate a safe "points-transfer" between two teams with built-in validation:

- Signature:** takes three arguments — the source team ID, the destination team ID, and the number of points to move.
- Guard clause:** immediately checks IF pts < 0 and, if so, logs "Cannot transfer negative points" into your ErrorLog table and raises an exception to abort the rest of the operation.
- It then performs two UPDATE statements—debiting points from the from_team and crediting them to the to_team. Any other error raised during these updates is caught by the EXCEPTION WHEN OTHERS block, which logs the database error message into ErrorLog and re-raises the exception so the entire transaction rolls back.

```
CREATE OR REPLACE PROCEDURE transfer_team_points(
  from_team INT,
  to_team   INT,
  pts       NUMERIC
)
LANGUAGE plpgsql
AS $$$
BEGIN
  -- Guard against bad input
  IF pts < 0 THEN
    INSERT INTO ErrorLog(err_msg)
    VALUES ('Cannot transfer negative points');
    RAISE EXCEPTION 'Cannot transfer negative points';
  END IF;

  -- Perform the transfer
  UPDATE TeamStandings
    SET total_points = total_points - pts
  WHERE team_id = from_team;

  UPDATE TeamStandings
    SET total_points = total_points + pts
  WHERE team_id = to_team;
EXCEPTION WHEN OTHERS THEN
  -- Log any unexpected error, then re-raise it so the whole transaction aborts
  INSERT INTO ErrorLog(err_msg) VALUES (SQLERRM);
  RAISE;
END;
$$;
```

After this, we attempt to insert two rows for season 1 (team 1 with 100 points & rank 1, and team 2 with 100 points & rank 2). If a row with the same (season_id, team_id) already exists, it instead **updates** that row's total_points and rank to the new values.

```
Query  Query History
1 ✓ INSERT INTO TeamStandings (season_id, team_id, total_points, rank)
2   VALUES (1,1,100,1), (1,2,100,2)
3   ON CONFLICT (season_id,team_id) DO UPDATE
4     SET total_points = EXCLUDED.total_points,
5         rank          = EXCLUDED.rank;
6 |
```

Data Output Messages Notifications

INSERT 0 2

Query returned successfully in 101 msec.

Now we invoke the `transfer_team_points(1, 2, 10)` procedure. It successfully debits 10 points from team 1 and credits 10 points to team 2 in one atomic operation (with all the same error-checking and logging built into the procedure). Since 10 is nonnegative and both updates succeed, we see a clean “CALL” message and no errors.

```
Query Query History
1 CALL transfer_team_points(1, 2, 10);
2
```

Data Output Messages Notifications

CALL

Query returned successfully in 62 msec.

The below screenshot shows the post-transfer standings for season 1, limited to teams 1 and 2. The SELECT returns each row’s primary key (`standing_id`), the `season_id`, the `team_id`, and their current `total_points` and `rank`, so we can verify that `transfer_team_points` call correctly updated each team’s points and left their ranks intact.

```
Query Query History
1 SELECT standing_id, season_id, team_id, total_points, rank
2   FROM TeamStandings
3 WHERE season_id = 1 AND team_id IN (1,2);
4
```

Data Output Messages Notifications

standing_id	season_id	team_id	total_points	rank
3001	1	1	100.00	1
3002	1	2	100.00	2

The below screenshot shows what happens when we call the procedure with a negative amount: the guard fires, the error message is inserted into ErrorLog (notice messages), and then an exception is raised (“Cannot transfer negative points”), aborting the transaction.

```
Query Query History
1 CALL transfer_team_points(2, 1, -5);
2
```

Data Output Messages Notifications

NOTICE: Error logged at 2025-04-27 17:53:26.020378: Cannot transfer negative points
NOTICE: Error logged at 2025-04-27 17:53:26.020378: Cannot transfer negative points

ERROR: Cannot transfer negative points
CONTEXT: PL/pgSQL function transfer_team_points(integer,integer,numeric) line 7 at RAISE

SQL state: P0001

The below final query pulls the three most recent rows from ErrorLog table—here showing that the last logged error was “cannot commit while a subtransaction is active,” along with the exact timestamp it was recorded.

Query Query History

```
1 SELECT *
2   FROM ErrorLog
3 ORDER BY log_id DESC
4 LIMIT 3;
```

Data Output Messages Notifications

log_id	err_time	err_msg
1	2025-04-27 17:18:09.765004	cannot commit while a subtransaction is active

What Happens on Transaction Abort?

- PL/pgSQL subtransactions.** The procedure body runs in a subtransaction.
- Guard-clause exception** aborts that subtransaction, rolling back anything done up to (but not including) the EXCEPTION block.
- EXCEPTION handler’s INSERT** into ErrorLog executes in the outer transaction context and thus survives.
- Reraising** the exception then aborts the outer transaction, rolling back any other changes (e.g. an explicit BEGIN; ... COMMIT; block is aborted).

Justification: by catching failures in a subtransaction, logging them, and then re-throwing, we ensure that no partial updates to business tables slip through, while still preserving an audit trail of the failure itself.

X. TASK 7 – INDEXING AND QUERY EXECUTION ANALYSIS

Below are 3 real world examples in the database where the planner falls back to expensive sequential along with simple B-tree indexing fixes that significantly cut I/O and CPU:

1. Aggregating every driver’s total points

```
Query Query History
1 EXPLAIN ANALYZE
2 SELECT driver_id, SUM(points_awarded)
3   FROM RaceResults
4 GROUP BY driver_id;
5
```

Data Output Messages Notifications

QUERY PLAN					
text					
1	HashAggregate (cost=68.00..91.46 rows=1878 width=36) (actual time=3.341..3.972 rows=1878 loops=1)				
2	Group Key: driver_id				
3	Batches: 1 Memory Usage: 881kB				
4	> Seq Scan on raceresults (cost=0.00..53.00 rows=3000 width=10) (actual time=0.084..0.673 rows=3001 loop=1)				
5	Planning Time: 0.196 ms				
6	Execution Time: 5.202 ms				

Initial Plan:

HashAggregate over Seq Scan on the full RaceResults table

Cost: ~ 68...91

Execution: ~ 5.2 ms to read 3000 rows and build the hash.

Because there’s no index on `driver_id`, PostgreSQL has to touch every row in the heap and build an in-memory hash table.

Improvement

```
1 -- Speed up GROUP BY on RaceResults(driver_id)
2 CREATE INDEX idx_rr_driver ON RaceResults(driver_id);
```

A B-tree on `driver_id` lets the planner consider an **Index Only Scan** (or a much smaller working set in cache) instead of a wide heap scan. Even if it still chooses a hash-aggregate, overall wall-clock time typically drops ~40–50% (the test fell from ~5.2 ms → ~2.6 ms).

```

1 ✓ EXPLAIN ANALYZE
2 SELECT driver_id, SUM(points_awarded)
3   FROM RaceResults
4 GROUP BY driver_id;

```

Query Plan text

- 1 HashAggregate (cost=68.02..91.49 rows=1878 width=36) (actual time=1.701..2.319 rows=1878 loops=1)
 Group Key: driver_id
 Batches: 1 Memory Usage: 881kB
 -> Seq Scan on raceresults (cost=0.00..53.01 rows=3001 width=10) (actual time=0.021..0.265 rows=3001 loops=1)
 Planning Time: 2.571 ms
 Execution Time: 2.620 ms

2. Give me the 5 most recent races for driver 1234

```

1 ✓ EXPLAIN ANALYZE
2 SELECT *
3   FROM RaceResults
4 WHERE driver_id = 1234
5 ORDER BY race_id DESC
6 LIMIT 5;

```

Query Plan text

- 1 Limit (cost=60.51..60.52 rows=1 width=30) (actual time=0.276..0.277 rows=1 loops=1)
 -> Sort (cost=60.51..60.52 rows=1 width=30) (actual time=0.275..0.275 rows=1 loops=1)
 Sort Key: race_id DESC
 Sort Method: quicksort Memory: 25kB
 -> Seq Scan on raceresults (cost=0.00..60.50 rows=1 width=30) (actual time=0.246..0.266 rows=1 loops=1)
 Filter: (driver_id = 1234)
 Rows Removed by Filter: 3000
 Planning Time: 0.232 ms
 Execution Time: 0.310 ms

Initial Plan:

Seq Scan on 3000 rows to filter driver_id = 1234
quicksort of the matching rows

Cost: ~ 60...60.5

Execution: ~ 0.31 ms

Even though it only returns 5 rows, PostgreSQL must scan every row in the table, then sort the few matches.

Improvement:

```
-- Speed up recent-race lookup (filter + sort)
CREATE INDEX idx_rr_driver_raceid ON RaceResults(driver_id, race_id DESC);
```

A composite index on (driver_id, race_id DESC) gives an Index Scan ... ORDER BY ... LIMIT plan, which drastically drops cost to ~ 0.28...8.30 and execution to ~ 0.11 ms, since the top 5 lives right in the index leaf pages.

```

1 ✓ EXPLAIN ANALYZE
2 SELECT *
3   FROM RaceResults
4 WHERE driver_id = 1234
5 ORDER BY race_id DESC
6 LIMIT 5;

```

Query Plan text

- 1 Limit (cost=0.28..8.30 rows=1 width=30) (actual time=0.093..0.094 rows=1 loops=1)
 -> Index Scan using idx_rr_driver_raceid on raceresults (cost=0.28..8.30 rows=1 width=30) (actual time=0.091..0.092 rows=1 loops=1)
 Index Cond: (driver_id = 1234)
 Planning Time: 0.148 ms
 Execution Time: 0.111 ms

3. Joining recent races to their circuits

```

1 ✓ EXPLAIN ANALYZE
2 SELECT r.race_name, c.circuit_name
3   FROM Races r
4 JOIN Circuits c ON r.circuit_id = c.circuit_id
5 WHERE r.race_date > '2024-01-01';

```

Query Plan text

- 1 Hash Join (cost=97.50..166.10 rows=801 width=45) (actual time=2.638..3.383 rows=807 loops=1)
 Hash Cond: (r.circuit_id = c.circuit_id)
 -> Seq Scan on races r (cost=0.00..66.50 rows=801 width=28) (actual time=0.048..0.515 rows=807 loops=1)
 Filter: (race_date > '2024-01-01')
 Rows Removed by Filter: 2194
 -> Hash (cost=60.00..60.00 rows=3000 width=25) (actual time=2.568..2.569 rows=3000 loops=1)
 Batches: 4096 Batches: 1 Memory Usage: 204kB
 -> Seq Scan on circuits c (cost=0.00..60.00 rows=3000 width=25) (actual time=0.043..1.649 rows=3000 loops=1)
 Planning Time: 1.849 ms
 Execution Time: 3.460 ms

Initial Plan:

Seq Scan on Races (filter by date)

Seq Scan on Circuits

Hash Join on circuit_id

Cost: ~ 97...166

Execution: ~ 3.46 ms

Without an index on race_date, all 3000 rows of Races must be examined, then joined to every row of Circuits (also 3000 rows).

Improvement:

```
-- Speed up date-range queries on Races
CREATE INDEX idx_races_date ON Races(race_date);
```

This lets the planner perform a **Bitmap Index Scan** on Races to prefilter only ~ 800 qualifying rows, cutting down the hash-build side and reducing ~ 10–20% off execution time. The PK on circuit_id already gives us an index on the other side of the join.

XI. CONCLUSION

Our project has delivered a fully-featured, end-to-end Formula 1 Racing Season Management Database in PostgreSQL, beginning with a carefully normalized schema of ten relations (Seasons, Drivers, Teams, Circuits, Races, RaceResults, QualifyingResults, PitStops, Penalties, TeamStandings) populated with over 3000 synthetic records each via a Python-Faker data-generation script. We enforced data integrity through primary and foreign keys, UNIQUE and CHECK constraints, and verified Boyce–Codd Normal Form across all tables. To demonstrate real-world utility, we executed more than ten SQL operations—ranging from inserts, updates, and deletes to complex analytical SELECTs with JOINs, GROUP BY, subqueries, and window functions and wrapped common workflows in reusable stored procedures and functions (e.g. adding race results with duplicate checks, updating driver nationalities, cascading deletes). We further improved the system by building a transactional “points transfer” procedure with in-database error logging and an AFTER-INSERT trigger on an ErrorLog table, ensuring that failures roll back cleanly and are automatically recorded. Finally, we used EXPLAIN ANALYZE to pinpoint three costly queries and applied targeted B-tree indexes to transform them into efficient index scans, proving the database’s scalability as data volumes grow. Overall, our project delivers a robust, maintainable, and high-performance platform for both real-time updates and deep analytical queries—far surpassing the flexibility and reliability of traditional spreadsheet solutions in the fast-paced world of Formula 1 racing.

REFERENCES

- [1] Faker Documentation. (2025). *Faker: Python Library for Generating Fake Data*. <https://faker.readthedocs.io/>
- [2] Formula 1 Official Website. (2025). <https://www.formula1.com/>
- [3] Formula 1 Wikipedia. https://en.wikipedia.org/wiki/Formula_One
- [4] PostgreSQL Global Development Group. (2025). *PostgreSQL 15.1 Documentation*. <https://www.postgresql.org/docs/>
- [5] PostgreSQL tutorial <https://neon.tech/postgresql/tutorial>
- [6] Transactions <https://www.postgresql.org/docs/current/tutorial-transactions.html>
- [7] Indexing <https://www.postgresql.org/docs/current/indexes.html>
- [8] Indexing <https://use-the-index-luke.com/>

Phase 1

Sharvari Mayekar	Vishnu K Menon
Defining relations/ attributes – 20 pts	E/R diagram - 20 pts
Target user – 20 pts	BCNF, Data acquisition – 20 pts
Problem statement – 10 pts	Problem statement – 10 pts
Total : 50 pts	Total : 50 pts

Phase 2

Sharvari Mayekar	Vishnu K Menon
Project Report – 10 points	Task 6 Transaction and Triggers- 25 points
Task 5 SQL Queries – 25 points	Presentation and Demo – 10 points
Demo – 5 points	Task 7 - Indexing and Query Execution Analysis- 15 points
Task 7 -Indexing and Query Execution Analysis- 10 points	
Total : 50 points	Total : 50 points